

Project: Object Recognition in Photographs

Changning Shou, Jun Yong Go

Agenda

- Simple CNN for object recognition
 - Deeper CNN
 - Even deeper CNN
 - Experimenting with data augmentation
 - Other Techniques
-

Setup

- Nvidia GTX 1060
- CNN
 - Better than ANN's
 - Translational invariance
 - Number of pooling layers and filter size are secondary to data augmentation ([Source](#))
 - Faster due to parameter sharing
 - Each unit in a feature map shares the same weight matrix ([Source](#))
- CIFAR-10 Dataset
- Keras

Photograph Object Recognition Dataset

CIFAR-10 Dataset

- 60,000 photos divided into 10 classes
 - 50,000 / 10,000
- 32 x 32 pixel images
- Performances:
 - 94% Human performance
 - 96% State-of-the-art results

airplane



automobile



bird



cat



deer



dog



frog



horse



ship



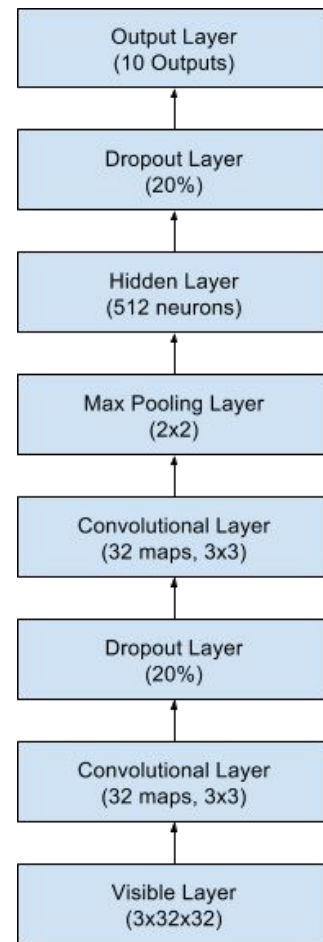
truck



[Source](#)

Simple CNN: Architecture

- Pattern:
 - Convolutional
 - Dropout
 - Convolutional
 - Max Pooling



[Source](#)

Simple CNN: Parameters

- Seed: 7
- Momentum: 0.9
- Learning Rate: 0.01
- Epochs: 25
- Batch Size: 32
- Logarithmic loss function

71%

Simple CNN Accuracy

Deep CNN: Architecture

- Same pattern as in simple CNN
- Pattern repeated 3 times with 32, 64, and 128 feature maps
 - Increasing number of feature maps with decreasing size given the max pooling layers
 - Additional and larger Dense layers at the output end to better translate the large number of feature maps to class values

Deep CNN: Parameters

- Seed: 7
- Momentum: 0.9
- Learning Rate: 0.01
- Epochs: 25
- Batch Size: 64
- Logarithmic loss function

77%

Deep CNN Accuracy

~15 min to run

Deeper CNN - Feature Maps: Architecture

- More feature maps closer to input
 - Deep CNN → 6 convolutional layers
 - 32, 32, 64, 64, 128, 128 → 64, 64, 64, 128, 128, 128

Deeper CNN - Feature Maps: Parameters

- Seed: 7
- Momentum: 0.9
- Learning Rate: 0.01
- Epochs: 25
- Batch Size: 64
- Logarithmic loss function

80%

Deeper CNN Accuracy - More Feature Maps

Deeper CNN - Less Pooling: Architecture

- Less aggressive pooling
- $2 \times 2 \rightarrow 3 \times 3$

Deeper CNN - Less Pooling: Parameters

- Seed: 7
- Momentum: 0.9
- Learning Rate: 0.01
- Epochs: 25
- Batch Size: 64
- Logarithmic loss function

77%

Deeper CNN Accuracy - Less Aggressive Pooling

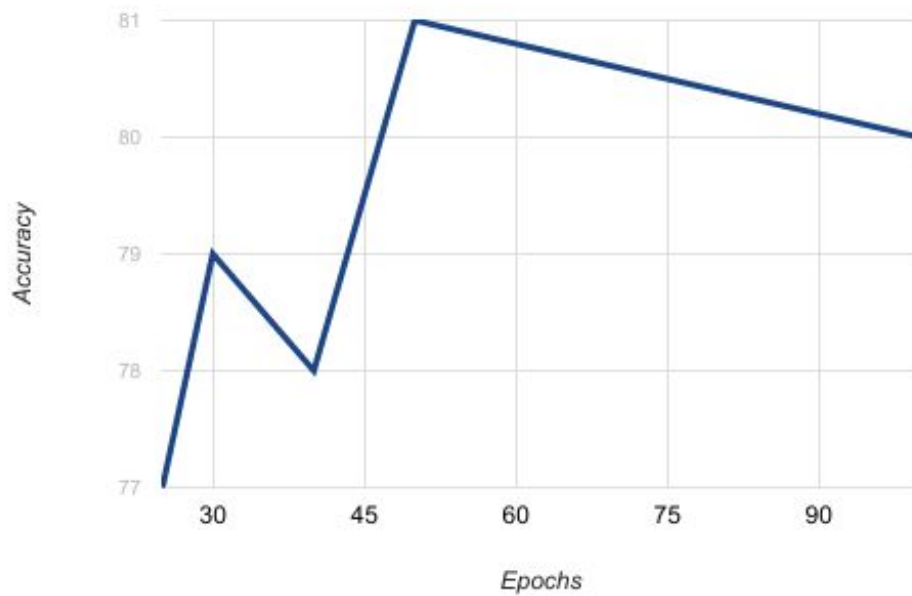
82%

Deeper CNN Accuracy -
More Features and Less Aggressive Pooling

Experiment: Epochs

- Deep CNN architecture
- Epochs: 25, 30, 40, 50, 100

Accuracy of Deep CNN with Different Epochs



Experiment: Data Augmentation

- Rotation: 180 degrees
- Flipping: horizontal and vertical

59%

Deeper CNN Accuracy - Rotations/Flips - 25 Epochs

66%


Deeper CNN Accuracy - Rotations/Flips - 100 Epochs

Other Techniques


- Fractional Max Pooling - Benjamin Graham
 - Pooling through overlapping squares or disjoint collection of rectangles
 - Other additions: translations, rotations, reflections, stretching...
 - Idea is to reduce spatial size of image by a factor of α ($1 < \alpha < 2$)
 - Alternative to (2×2) max pooling, stochastic pooling, which reduce size of hidden layers by factor of two
 - Intuition: More pooling layers and more opportunity to view and recognize features in the image
 - Introducing randomness to pooling
 - Contrast to stochastic pooling, which has randomness in its pooling
 - Random vs. Pseudorandom pooling
 - Performance: 96.53% with data augmentation

Other Techniques

- The All Convolutional Net - Jost Tobias Springenberg , Alexey Dosovitskiy, Thomas Brox, Martin Riedmiller
 - Replace max pooling with convolutional layer with increased stride (strided convolution / deconvolutional layer / transposed convolutional layer)
 - Homogeneous network consisting of convolutional layers
 - Results give evidence to effectiveness of small convolutional layers and raises questions about the necessity of pooling in CNNs
 - Pooling layer is just a feature-wise convolution that uses average, p-norm, max function
 - Increasing stride of previous convolutional layer vs. Adding a convolutional layer with corresponding stride
 - Performance: 92% without data augmentation



Zero-padding, Stride = 1



Zero-padding, Stride = 2

[Source](#)